


Problem Set 9

Daniel Wang (S01435533)

1. The maximum number of edges in a graph is when the graph is complete. In this scenario, $|E| = O(|V|^2)$, so $O(|E|)$ is not necessarily $O(|V|)$. However, we can induce that $O(\log|E|) = O(\log|V|^2) = O(2 \cdot \log|V|) = O(\log|V|)$
2. The graph of courses can be visualized as follow:




We can use topological sort based on DFS to solve this question. A possible answer can be derived from the following steps, where courses in boldface represents that they are being appended to the current end of list:

- (1) LA15 → LA31 → LA32 → **LA169**, backtrack to LA32
- (2) LA32 → **LA126**, backtrack to **LA32**, **LA31**, LA15
- (3) LA15 → LA16 → **LA141**, backtrack to LA16
- (4) LA16 → **LA127**, backtrack to **LA16**, **LA15**
- (5) LA22, adjacent nodes are visited, so append **LA22**

Based on previous traversal sequences, we can derive a viable course plan using a reverse order:

LA22 → LA15 → LA16 → LA127 → LA141 → LA31 → LA32 → LA126 → LA169

3. The adjacency list can be translated into the following graph:



The DFS sequence can be derived as follows:

- (1) $0 \rightarrow 1$, visited = {0}
- (2) $1 \rightarrow 2$, visited = {0, 1}
- (3) $2 \rightarrow 3$, visited = {0, 1, 2}
- (4) $3 \rightarrow 5$, visited = {0, 1, 2, 3}
- (5) $5 \rightarrow 4$, visited = {0, 1, 2, 3, 5}
- (6) $4 \rightarrow 6$, visited = {0, 1, 2, 3, 4, 5}
- (7) $6 \rightarrow 7$, visited = {0, 1, 2, 3, 4, 5, 6}
- (8) 7 halts since all neighbors are being visited, visited = {0, 1, 2, 3, 4, 5, 6, 7}
- (9) Since all nodes are being visited, DFS halts

Thus, the DFS sequence is $0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 5 \rightarrow 4 \rightarrow 6 \rightarrow 7$

4. The BFS sequence can be derived as follows:

- (1) Initialize queue = [0], mark = {0}
- (2) Pop 0, queue = [1, 2, 3], mark = {0, 1, 2, 3}
- (3) Pop 1, queue = [2, 3], mark = {0, 1, 2, 3}
- (4) Pop 2, queue = [3], mark = {0, 1, 2, 3}
- (5) Pop 3, queue = [5], mark = {0, 1, 2, 3, 5}
- (6) Pop 5, queue = [4, 6], mark = {0, 1, 2, 3, 4, 5, 6}
- (7) Pop 4, queue = [6, 7], mark = {0, 1, 2, 3, 4, 5, 6, 7}
- (8) Pop 6, queue = [7], mark = {0, 1, 2, 3, 4, 5, 6, 7}
- (9) Pop 7, queue = [], mark = {0, 1, 2, 3, 4, 5, 6, 7}
- (10) Since the queue is empty, BFS halts

Thus, the BFS sequence is $0 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 5 \rightarrow 4 \rightarrow 6 \rightarrow 7$

5. Here I assume that for each element in the heap is a pair, where the first index represents the priority and the second one means the value. I will use a hash map when constructing the heap, where the key is the item name, and the value is the index inside current heap. When heapify, I will not only swap the element in the array but also change the record to match current swapped status. The Python code is shown below:

```
def heapify(arr, start, end, name_to_idx):
    """ min-heapify the array within a range [start, end],
    update name_to_idx"""
    for i in range(end//2-1, start-1, -1):
        l, r = 2*i+1, 2*i+2
        max_id = i
        if arr[max_id][0] > arr[l][0]:
            max_id = l
        if r < end and arr[max_id][0] > arr[r][0]:
            max_id = r

        # swap name_to_idx
        name1, name2 = arr[i][1], arr[max_id][1]
        name_to_idx[name1] = max_id
        name_to_idx[name2] = i

        # swap priority
        arr[i], arr[max_id] = arr[max_id], arr[i]
        if max_id != i:
            heapify(arr, max_id, end, name_to_idx)

def update_priority(arr, name, val, name_to_idx):
    idx = name_to_idx[name]
    arr[idx][0] = val
    heapify(arr, 0, len(arr), name_to_idx)
```

When we run the following code in the main loop where we try to change the priority of A from 1 to 6:

```
arr = [[1, "A"], [0, "B"], [3, "C"], [2, "D"], [5, "E"]]
n = len(arr)
name_to_idx = {arr[i][1]: i for i in range(n)}
heapify(arr, 0, n, name_to_idx)

# before update
print(arr)
print(name_to_idx)
print("-----")

# after update
update_priority(arr, "A", 6, name_to_idx)
print(arr)
print(name_to_idx)
```

The console output satisfies the property of heap and is demonstrated below:

```
[[0, 'B'], [1, 'A'], [3, 'C'], [2, 'D'], [5, 'E']]
```

```
{'A': 1, 'B': 0, 'C': 2, 'D': 3, 'E': 4}
```

```
[[0, 'B'], [2, 'D'], [3, 'C'], [6, 'A'], [5, 'E']]
```

```
{'A': 3, 'B': 0, 'C': 2, 'D': 1, 'E': 4}
```

6. Assume the input of Dijkstra algorithm contains *graph*, *weights*, and *src*. Where *graph* is stored as an adjacency list, *weights* is a dictionary where key is the edge and value is the path cost, and *src* is the vertex name of source node. The modified code is shown as follows, where we adopt the same function from Q.5:

```
def dijkstra(graph, weights, src):
    all_nodes = set(graph.keys())

    pq = [] # (vertex, current distance)
    dists = {} # {vertex: distance}
    name_to_idx = {} # {vertex: current index in pq}
    for (i, node) in enumerate(all_nodes):
        pq.append([float("inf"), node])
        dists[node] = float("inf")
        name_to_idx[node] = i

    # initialize for source node
    dists[src] = 0
    update_priority(pq, src, 0, name_to_idx)


    # traverse all nodes
    result = []
    while len(pq):
        # pop smallest while maintain name_to_idx
        dist, node = pq[0]
        last_name = pq[-1][1]
        pq[0] = pq[-1]
        name_to_idx[last_name] = 0

        pq.pop()
        heapify(pq, 0, len(pq), name_to_idx)

        # add result and traverse its neighbor
        result.append(node)
        for node_next in graph[node]:
            new = dist[node] + weights[(node, node_next)]
            if new < dists[node_next]:
                dists[node_next] = new
                update_priority(pq, node_next, new, name_to_idx)


    return result
```

7. Assume the following acyclic graph with a source node of A:




The minimum spanning tree should be $\{AC, AB, BD\}$, having a cost of -96. Using Dijkstra started with node A, it traverses to C and then to D, backtracking to C and A, and traverses to B. At this time, D is already being marked beforehand, so the edge BD will be ignored. In this scenario, Dijkstra yields a suboptimal result $\{AC, AD, AB\}$, yielding a cost of 5 and is greater than -96.

8. Consider the following acyclic graph and assume the source node is A:



The optimal solution is $\{AB, AD, BC, CE\}$, yielding a cost of 0. Now we offset the value of edges by +2 so that all edges become positive. It becomes:



Using Dijkstra started from A, the appended edges will have the order of AB, AD, BC, DE. The edge CE will not be considered since the path $A \rightarrow D \rightarrow E$ has a cost of 6, while the path $A \rightarrow B \rightarrow C \rightarrow E$ has a cost of 7. In this scenario, Dijkstra yields a result $\{AB, AD, BC, DE\}$, yielding a cost of 2 and is greater than 0.

9. The method can be decomposed into 2 steps.

- (1) Since universal sink node can have at most 1, the first steps we can try is to eliminate $n-1$ possible nodes by judging $A[i][j]$. If $A[i][j] = 0$, it means there is no link from node i to node j , implying that j cannot be a sink node. Thus, we can eliminate j from candidate pool. If $A[i][j] = 1$, it means there is a link from node i to node j , meaning that i cannot be a sink node. Doing $n-1$ iterations will make the number of candidates to be 1.
- (2) For the current candidate, we will search whether all nodes are connecting to it

and it cannot connect to others. The Python code is shown as follows:


```
def has_sink_node(A):
    n = len(A)

    # make i < j, and eventually j will be n
    # where i will be a possible candidate
    i, j = 0, 1
    while j < n:
        if A[i][j] == 0: # j impossible
            j = j + 1
        else: # i impossible
            i, j = j, j + 1

    # check whether i is really sink node
    idx = i
    for i in range(n):
        if i == idx:
            continue
        if not (A[i][idx] == 1 and A[idx][i] == 0):
            return False

    return True
```

10. The graph can be represented as follow:



Kruskal's algorithm tries to find smallest edge in the graph for each iteration. The algorithm stops when all nodes are being visited. Let assume *result* be the path list, and *conn* be the connected components for current iteration.

- (1) BE = 4 , result = [BE], conn = [{B,E}]
- (2) AD = 5, result = [BE, AD], conn = [{B,E}, {A,D}]
- (3) DF = 6, result = [BE, AD, DF], conn = [{B,E}, {A,D,F}]
- (4) AC = 7, result = [BE, AD, DF, AC], conn = [{B,E}, {A,C,D,F}]
- (5) BC = 8, result = [BE, AD, DF, AC, BC], conn = [{A,B,C,D,E,F}]
- (6) CE = 10, since C and E are in the same connected component, skip it
- (7) FG = 11, result = [BE, AD, DF, AC, BC, FG], conn = [{A,B,C,D,E,F, G}]
- (8) Since all nodes are in the same connected component, break the loop

Eventually, the minimum spanning tree will look like below:

